

Coordination of Robotic Networks: On Task Allocation and Vehicle Routing


Francesco Bullo



Center for Control,
Dynamical Systems & Computation
University of California at Santa Barbara
<http://motion.me.ucsb.edu>

Electrical and Computer Engineering Department
Carnegie Mellon University
April 30, 2009

Acknowledgements



Stephen L. Smith

SLS, FB: "Monotonic target assignment for robotic networks," *IEEE Trans Automatic Ctrl*, 54 (10), 2009

SLS, SDB, FB: "Finite-time pursuit of translating targets in a dynamic and stochastic environment," *Proc CDC*, Shanghai, 2009, submitted



Shaunak D. Bopardikar


SDB, SLS, FB, JH: "Dynamic vehicle routing for translating demands," *IEEE Trans Automatic Ctrl*, 2009, submitted

Applications of autonomous systems


- Unmanned vehicles
- Equipped with **suite of sensors**
- **Inaccessible** environments

Civilian applications:

- Environmental monitoring:
 - **Measure** weather systems
 - **Observe** animal species
 - **Detect** and **assess** wildfires
- Search and rescue missions
- Space exploration
- Monitoring infrastructure



Slocum glider



NASA – next generation Mars rover

Applications of autonomous systems

Military applications:

- Surveillance
- Reconnaissance missions
- Perimeter defense and security
- Expenditures of \$60 billion over next 10 years



Globalhawk



Aerovironment Wasp

The future of autonomy

Current missions (typical scenario):

- single vehicle or few decoupled vehicles
- pre-specified task
- tightly coupled with human control

Future missions

- 1 Fleets (swarms) of networked vehicles
- 2 Complex sets of tasks that evolve during execution
- 3 Increased autonomy, humans as supervisors

Requires real-time **task allocation** and **vehicle routing**


Task allocation

Given:

- a group of vehicles, and
- a set of tasks

Task example:

take a picture at a location



Task allocation

Decide which vehicles should perform which tasks.

- **Centralized**: operator assigns vehicles to tasks
(requires vehicle positions, workloads, etc.)
- **Distributed**: vehicles divide tasks among themselves


Task allocation

Given:

- a group of vehicles, and
- a set of tasks

Task example:

take a picture at a location



Task allocation


Decide which vehicles should perform which tasks.

- **Centralized:** operator assigns vehicles to tasks
(requires vehicle positions, workloads, etc.)
- **Distributed:** vehicles divide tasks among themselves

Vehicle routing

Given:

An allocation of tasks to vehicles



Vehicle routing


Determine a path that allows each vehicle to complete its tasks.

- Task A is of higher priority than task B
- A task requires multiple vehicles: vehicles need to rendezvous
- Task locations are not stationary

Vehicle routing

Given:

An allocation of tasks to vehicles



Vehicle routing


Determine a path that allows each vehicle to complete its tasks.

- Task A is of higher priority than task B
- A task requires multiple vehicles: vehicles need to rendezvous
- Task locations are not stationary

Vehicle routing

Given:

An allocation of tasks to vehicles



Vehicle routing


Determine a path that allows each vehicle to complete its tasks.

- Task *A* is of higher priority than task *B*
- A task requires multiple vehicles: vehicles need to rendezvous
- Task locations are not stationary

Vehicle routing

Given:

An allocation of tasks to vehicles



Vehicle routing


Determine a path that allows each vehicle to complete its tasks.

- Task A is of higher priority than task B
- A task requires multiple vehicles: vehicles need to rendezvous
- Task locations are not stationary

Vehicle routing

Given:

An allocation of tasks to vehicles



Vehicle routing

Determine a path that allows each vehicle to complete its tasks.

- Task A is of higher priority than task B
- A task requires multiple vehicles: vehicles need to rendezvous
- Task locations are not stationary


Dynamic and distributed aspects

Distributed:

Vehicles have only local information

Dynamic:

- Existing tasks evolve over time
- New tasks arise in real-time
- Number of vehicles changes



Complete solution **cannot be computed off-line**.

As new information becomes available, vehicles must


- re-allocate tasks
- re-plan paths

Technical approach:

structure, fundamental limits, efficient algorithms

For a distributed/dynamic problem:

- 1 Identify **underlying problem structure**
e.g., adimensional analysis, intrinsic regimes,
phase transitions in parameter space
- 2 Determine **fundamental limits** on performance
- 3 Design provably **efficient algorithms**



The remainder of the talk

Illustrate

problem structure, fundamental limits, efficient algorithms


via two scenarios:

- 1 Distributed Task Allocation
motivated by a surveillance application
- 2 Dynamic Vehicle Routing
motivated by a perimeter defense application

Distributed task allocation

A distributed task allocation problem

- n omnidirectional vehicles
 - limited comm. range and bandwidth
- $m \leq n$ task locations
 - once task is reached by a vehicle, vehicle is forever engaged



Two problem scenarios:


- 1 Supervisor broadcasts all task locations to each vehicle
- 2 Vehicles search for task locations with limited range sensor

Problem: distributed algorithm to

- allow group of vehicles to divide tasks among themselves
- minimize time until last task location is reached

A distributed task allocation problem

- n omnidirectional vehicles
 - limited comm. range and bandwidth
- $m \leq n$ task locations
 - once task is reached by a vehicle, vehicle is forever engaged



Two problem scenarios:


- 1 Supervisor broadcasts all task locations to each vehicle
- 2 Vehicles search for task locations with limited range sensor

Problem: distributed algorithm to

- allow group of vehicles to divide tasks among themselves
- minimize time until last task location is reached

A distributed task allocation problem

- n omnidirectional vehicles
 - limited comm. range and bandwidth
- $m \leq n$ task locations
 - once task is reached by a vehicle, vehicle is forever engaged



Two problem scenarios:

- 1 Supervisor broadcasts all task locations to each vehicle
- 2 Vehicles search for task locations with limited range sensor

Problem: distributed algorithm to

- allow group of vehicles to divide tasks among themselves
- minimize time until last task location is reached


Centralized solution

In the centralized setting, problem is **matching in a bipartite graph**

Specifically, **bottleneck matching**:
find a matching M which minimizes

$$\max_M d_{i,j}$$

Solvable in polynomial time



Distributed challenges

Multi-vehicle task allocation work:

- Auction based (Moore and Passino, 2007)
- Game theoretic (Arslan et al., 2007)
- Auction and consensus (Brunet, Choi and How, 2008)

Today, combination of **key challenges**:

- 1 **range constraint** and **lack of connectivity**
- 2 **tight bandwidth constraint**


and novel **goals**:

- 3 determine fundamental limits on scalability
- 4 develop provably efficient algorithms


Underlying structure: environment size regimes

If # of vehicles increases ($n \rightarrow +\infty$)

Then area $A(n)$ must increase to “make room”



Sparse: $A(n)/n \rightarrow +\infty$



Dense: $A(n)/n \rightarrow 0^+$

Critical: $A(n)/n \rightarrow \text{constant}$

Fundamental limits on completion time

Worst-case completion time

- # of tasks = # of vehicles ($m = n$)
- Broadcast or search scenario


	Sparse ($A(n) \gg n$)	Critical ($A(n) \approx n$)	Dense ($A(n) \ll n$)
Fundamental limit	$\Omega(\sqrt{nA(n)})$	$\Omega(n)$	$\Omega(A(n))$

Asymptotic notation: $T \in \Omega(n)$ implies there is $C > 0$ such that

T lower bounded by Cn

Two allocation algorithms


The Ring algorithm



- Compute **common ring**
- Broadcast scenario


Two allocation algorithms

The Ring algorithm



- Compute **common ring**
- Broadcast scenario

The Grid algorithm



- Elect **leader in each cell**
- Broadcast or search

Algorithms match fundamental limit

Worst-case time, ($\#$ of tasks m) = ($\#$ of vehicles n)

	Sparse	Critical	Dense
Fundamental limit	$\Omega(\sqrt{nA(n)})$	$\Omega(n)$	$\Omega(A(n))$
Ring Alg	$O(\sqrt{nA(n)})$	$O(n)$	$O(\sqrt{nA(n)})$
Grid Alg	$O(A(n))$	$O(n)$	$O(A(n))$

Efficient algorithms

- Ring Alg in sparse and critical environments
- Grid Alg in dense and critical environments


Additional stochastic results have been obtained

Summary of distributed task allocation

Distributed task allocation with communication constraints

The results:

- **problem structure**: sparse/critical/dense
- **fundamental limits** on completion time
- **efficient algorithms** in all three regimes



The technical approach utilizes:


- Distributed algorithms and networking
- Combinatorial optimization
- Random geometric graphs

Text: Distributed Control of Robotic Networks

Princeton Series in APPLIED MATHEMATICS

Distributed Control of Robotic Networks

A Mathematical Approach
to Motion Coordination Algorithms



Francesco Bullo
Jorge Cortés
Sonia Martínez

- 1 intro to distributed algorithms
(graph theory, synchronous networks,
and averaging algos)
- 2 geometric models and geometric
optimization problems
- 3 model for robotic, relative sensing
networks, and complexity
- 4 algorithms for rendezvous,
deployment, boundary estimation


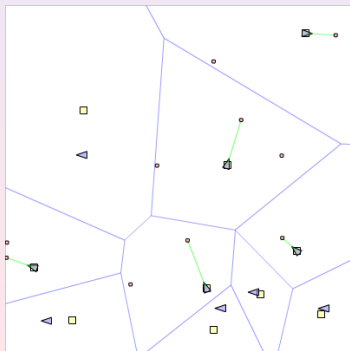
Status: Freely downloadable at
<http://coordinationbook.info>
with tutorial slides & software libraries.
Shortly on sale by Princeton Univ Press

Dynamic vehicle routing

Prior work on dynamic vehicle routing

Dynamic traveling repairperson problem

- Tasks arrive sequentially in time
- Each task location is randomly distributed in service region
- Each task requires **on-site service**



Key references

Key references


- Shortest path (Beardwood, Halton and Hammersly, 1959)
- Formulation on a graph (Psaraftis, 1988)
- Euclidean plane (Bertsimas and Van Ryzin, 1990–1993)

Recent developments in dynamic vehicle routing:

- Nonholonomic UAVs (Savla, Frazzoli, FB: TAC, (53)6 '08)
- Adaptation and decentralization (Pavone, Frazzoli, FB: TAC, sub '09)
- Distinct-priority targets (SLS, Pavone, FB, Frazzoli: SICON, sub '09)
- Heterogeneous vehicles and teaming (SLS, FB: SCL, sub '08)
- **Moving targets (SBD, SLS, FB: CDC & TAC, sub '09)**

A perimeter defense / boundary guarding problem

- Single vehicle with **unit speed**
- Task locations (targets):
 - **arrive sequentially** on a segment
 - **move vertically** with speed v
- Task completed if target **captured** before reaching deadline



Goal

Design policies that maximize expected fraction of targets captured

Assume that task arrivals are:

- Poisson in time with rate $\lambda \implies \mathbb{E}[N(\Delta t)] = \lambda \Delta t$
- uniformly distributed on line segment


Underlying problem structure

For fixed W , **problem parameters** are

- speed ratio v :

$$v = \frac{\text{target speed}}{\text{vehicle speed}}$$

- arrival rate λ
- deadline distance L



	$L = +\infty$ Stabilize queue	L is finite Maximize capture fraction
$v < 1$	translational path policy	translational path policy
$v \geq 1$	Not possible for any $\lambda > 0$	longest path policy


Underlying problem structure

For fixed W , **problem parameters** are

- speed ratio v :

$$v = \frac{\text{target speed}}{\text{vehicle speed}}$$

- arrival rate λ
- deadline distance L



	$L = +\infty$ Stabilize queue	L is finite Maximize capture fraction
$v < 1$	translational path policy	translational path policy
$v \geq 1$	Not possible for any $\lambda > 0$	longest path policy


Underlying problem structure

For fixed W , **problem parameters** are

- speed ratio v :

$$v = \frac{\text{target speed}}{\text{vehicle speed}}$$

- arrival rate λ
- deadline distance L




	$L = +\infty$ Stabilize queue	L is finite Maximize capture fraction
$v < 1$	translational path policy	translational path policy
$v \geq 1$	Not possible for any $\lambda > 0$	longest path policy

Fundamental limits for $L = +\infty$ and $\nu < 1$

For every policy:

$$\lambda \leq \frac{4}{\nu W}, \quad \text{for stability}$$





Fundamental limits for $L = +\infty$ and $\nu < 1$

For every policy:

$$\lambda \leq \frac{4}{\nu W}, \quad \text{for stability}$$

As $\nu \rightarrow 1^-$, for stability

$$\lambda \leq \frac{3\sqrt{2}}{W\sqrt{-\log(1-\nu)}}$$



Example of proof techniques

For stability, $\lambda \leq \frac{4}{vW}$

- 1 Distribution of unserved targets in region of area A :
 - Number is Poisson distributed with parameter $\lambda A/(vW)$
 - Conditioned on number, targets are uniform

- 2 Targets reachable in time T from (X, Y) are

$$\{(x, y) \mid (X - x)^2 + ((Y - vT) - y)^2 \leq T^2\}$$

- 3 Probability that closest target is not reachable in T seconds

$$\geq \exp(-\lambda \pi T^2/(vW))$$

- 4 Expected time to travel between targets

$$\mathbb{E}[\text{travel time}] \geq \frac{1}{2} \sqrt{\frac{vW}{\lambda}}$$

- 5 To capture all, $\lambda \mathbb{E}[\text{travel time}] \leq 1$

Translational path for $v < 1$


Shortest translational path policy

Input: Optimal location \mathbf{p}^*

- 1 If no targets, then move to \mathbf{p}^*
- 2 Else, capture all targets via **shortest translational path**
- 3 Repeat

Can compute \mathbf{p}^* to minimize:


- worst-case capture time
- expected capture time



speed ratio of 0.6


Stability of translational path for $v < 1$

Stability for $L = +\infty$




Stability of translational path for $v < 1$

Stability for $L = +\infty$




Stability of translational path for $v < 1$

Stability for $L = +\infty$



Maximize capture fraction for $v < 1$

Modify translational path policy



Fundamental limit

$$\text{cap fraction} \leq \min \left\{ 1, \frac{2}{\sqrt{v\lambda W}} \right\}$$

To analyze policy, assume

- speed ratio v is small
- arrival rate λ is large

Then, capture fraction

$$\geq \min \left\{ 1, \frac{1.4}{\sqrt{v\lambda W}} \right\}$$

Factor 1.42 of optimal

Numerical results suggest good performance away from limit

Where are we?

	$L = +\infty$ Stabilize queue	L is finite Maximize capture fraction
$\nu < 1$	translational path policy	modified trans. path policy
$\nu \geq 1$	Not possible for any $\lambda > 0$	

Where are we?


	$L = +\infty$ Stabilize queue	L is finite Maximize capture fraction
$\nu < 1$	translational path policy	modified trans. path policy
$\nu \geq 1$	Not possible for any $\lambda > 0$	

Where are we?

	$L = +\infty$ Stabilize queue	L is finite Maximize capture fraction
$\nu < 1$	translational path policy	modified trans. path policy
$\nu \geq 1$	Not possible for any $\lambda > 0$?


Maximize fraction of targets for $v \geq 1$

For $v \geq 1$, it is optimal to remain on deadline




Maximize fraction of targets for $v \geq 1$

For $v \geq 1$, it is optimal to remain on deadline




Maximize fraction of targets for $v \geq 1$

For $v \geq 1$, it is optimal to remain on deadline




Maximize fraction of targets for $v \geq 1$

For $v \geq 1$, it is optimal to remain on deadline



Maximize fraction of targets for $v \geq 1$

For $v \geq 1$, it is optimal to remain on deadline



- Reachability graph is directed and **acyclic**

Fundamental limit for $v \geq 1$

Noncausal information = *a priori* knowledge of
arrival time and location of every future target

Optimal performance with noncausal information

- 1 Compute infinite reachability graph of all future targets
- 2 Compute longest path in graph
- 3 Capture each target on path

Consequences for algorithm performance (capture fraction)

- noncausal performance can be computed
- noncausal performance is upper bound on causal performance


Capture fraction with $v \geq 1$: Longest path policy

Longest path (LP) policy

- 1 Compute the reachability graph of all unserved targets
- 2 Compute longest path in graph
- 3 Capture first target on path by intercepting on deadline
- 4 Repeat


Capture fraction for $L > vW$:

Factor $(1 - \frac{vW}{L})$ of optimal




Numerical capture fraction for $\nu \geq 1$

Environment with $W = 2$ and $L = 5$.




$\nu = 2$ and thus $L > \nu W$




$\nu = 5$ and thus $L < \nu W$

Numerical capture fraction for $\nu \geq 1$

Environment with $W = 2$ and $L = 5$.



$\nu = 2$ and thus $L > \nu W$



$\nu = 5$ and thus $L < \nu W$

Summary of boundary guarding

The results:

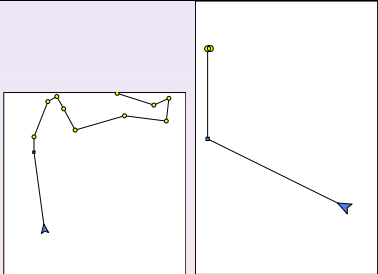
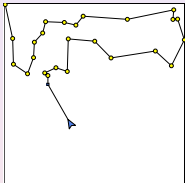
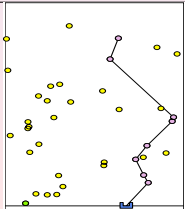
- Identified **four regimes**
- Derived **fundamental limits** on capture fraction
- Developed provably **efficient algorithms**

	$L = +\infty$ Stabilize queue	L is finite Maximize capture fraction
$\nu < 1$	translational path policy	translational path policy
$\nu \geq 1$	Not possible for any $\lambda > 0$	longest path policy

The technical approach utilizes:

- Stochastic processes and queueing
- Combinatorial optimization

Summary of boundary guarding: policies

	$L = +\infty$ Stabilize queue	L is finite Maximize capture fraction
$v < 1$		
$v \geq 1$	<p>Not possible for any $\lambda > 0$</p>	

Future autonomous missions

- Fleets (swarms) of networked vehicles
- Complex sets of tasks that evolve during execution
- Increased autonomy, humans as supervisors

Enabling technology: real-time **task allocation** and **vehicle routing**

Technical approach: Fundamental theory and algorithms

- 1 underlying **problem structure**
- 2 **fundamental limits** on performance
- 3 simple, provably **efficient algorithms**